

Generative Adversarial Networks (GANs)

A gentle Introduction

<https://github.com/toelt-llc/GANs-TensorFlow>

Umberto Michelucci

umberto.michelucci@toelt.ai

+41 79 396 7406

Who am I

- Head of AI Center of Excellence @ Helsana Versicherung AG
- Master in Theoretical Physics and PhD in Computer Science and Machine Learning
- Author of «Applied Deep Learning –A Case-Based Approach to Understanding Deep Neural Networks» (APRESS 2018)
- Author of «Advanced Applied Deep Learning – Convolutional Neural Networks and Object Detection» (APRESS 2019)
- Founder Toelt GmbH
- Google Developer Expert in Machine Learning
- More than 25 published papers in the last three years in Machine Learning

Helsana






Painter: she
would like to
learn to learn
to paint as Van
Gogh



Critic: she
would like to
learn be able
to tell real Van
Gogh
paintings from
fakes apart.

Step 1



Step 2 – The critic must become better too




TRUE



FAKE

The critic trains with
real Van Gogh
paintings



The critic trains with
clearly fake Van
Gogh paintings




Critic: she
would like to
learn be able
to tell real Van
Gogh
paintings from
fakes apart.



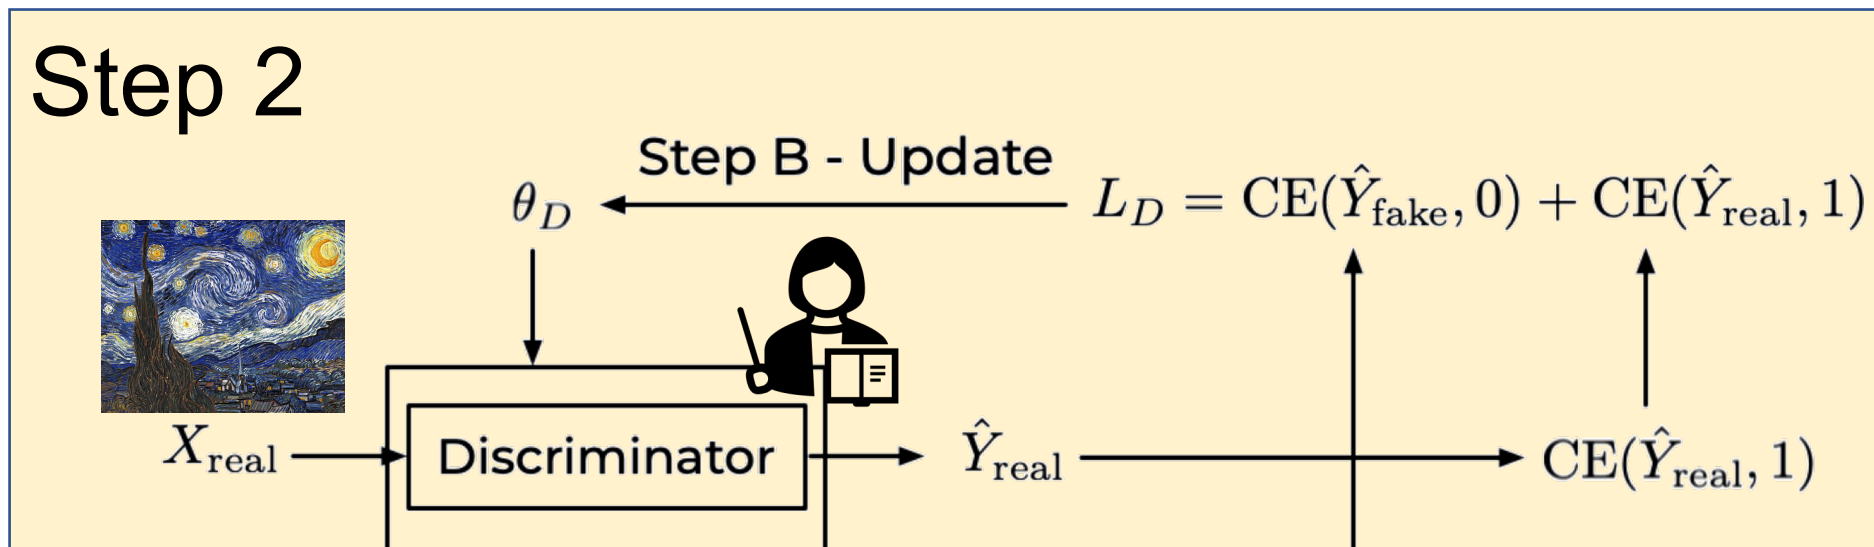
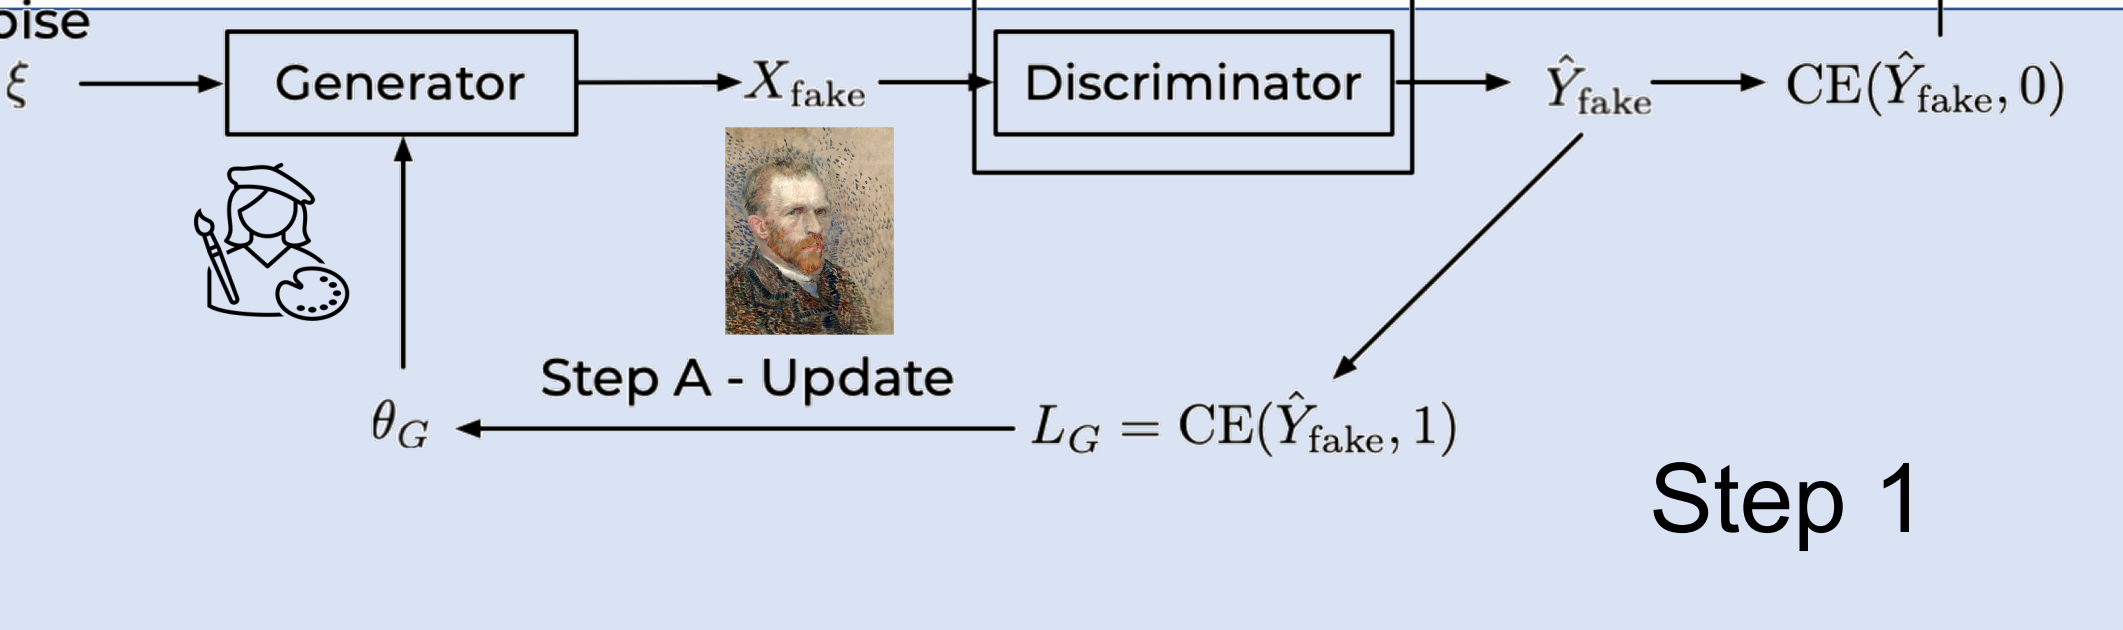
The critic becomes
slowly better

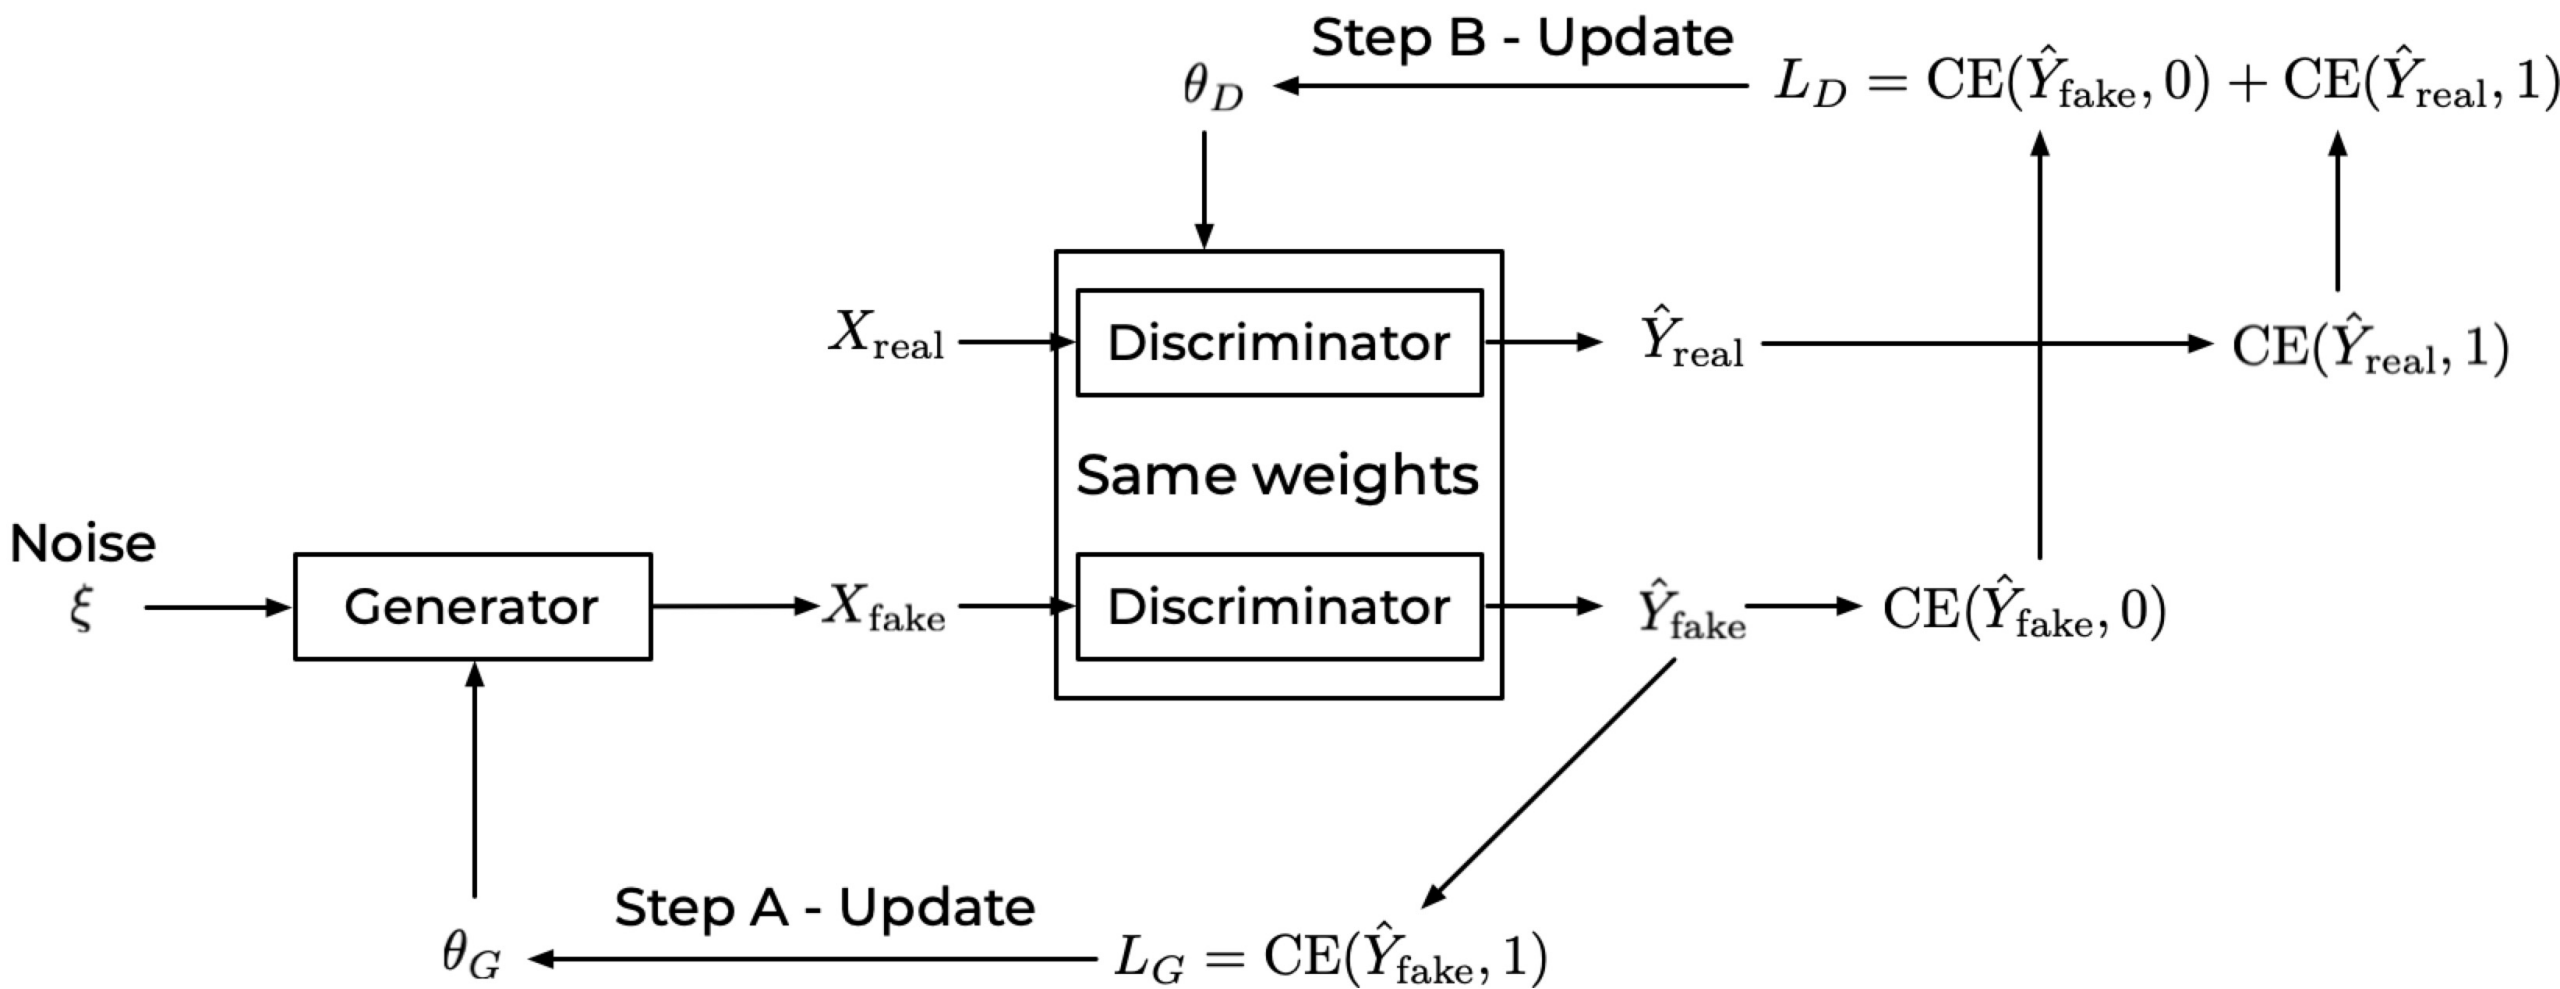
Reference – Loss Function (Cross Entropy, CE)



Ideas of the painter
(today I will paint a portrait,
a landscape, etc.)

Noise





How to use Keras for GANs

The most important aspect of using Keras for GANs is the necessity of building a custom training loop (and not using **compile()/fit()** approach).

```
def train_step(images):
    noise = tf.random.normal([BATCH_SIZE, noise_dim])

    with tf.GradientTape() as gen_tape, tf.GradientTape() as disc_tape:
        # Calculation of  $X_{fake}$ 
        generated_images = generator(noise, training=True) ←  $X_{fake}$ 

        # Calculation of  $\hat{Y}_{real}$ 
        real_output = discriminator(images, training=True) ←  $X_{real}$ 

        # Calculation of  $\hat{Y}_{fake}$ 
        fake_output = discriminator(generated_images, training=True)

        # Calculation of  $L_G$ 
        gen_loss = generator_loss(fake_output) ←  $L_G = CE(\hat{Y}_{fake}, 1)$ 

        # Calculation of  $L_D$ 
        disc_loss = discriminator_loss(real_output, fake_output) ←  $L_D = CE(\hat{Y}_{fake}, 0) + CE(\hat{Y}_{real}, 1)$ 

    # Gradients Calculation
    # Calculation of the gradients of  $L_G$  for backpropagation
    gradients_of_generator = gen_tape.gradient(gen_loss, generator.trainable_variables)

    # Calculation of the gradients of  $L_D$  for backpropagation
    gradients_of_discriminator = disc_tape.gradient(disc_loss, discriminator.trainable_variables)

    # Training Steps A and B
    # Step A
    generator_optimizer.apply_gradients(zip(gradients_of_generator, generator.trainable_variables))

    # Step B
    discriminator_optimizer.apply_gradients(zip(gradients_of_discriminator, discriminator.trainable_variables))
```

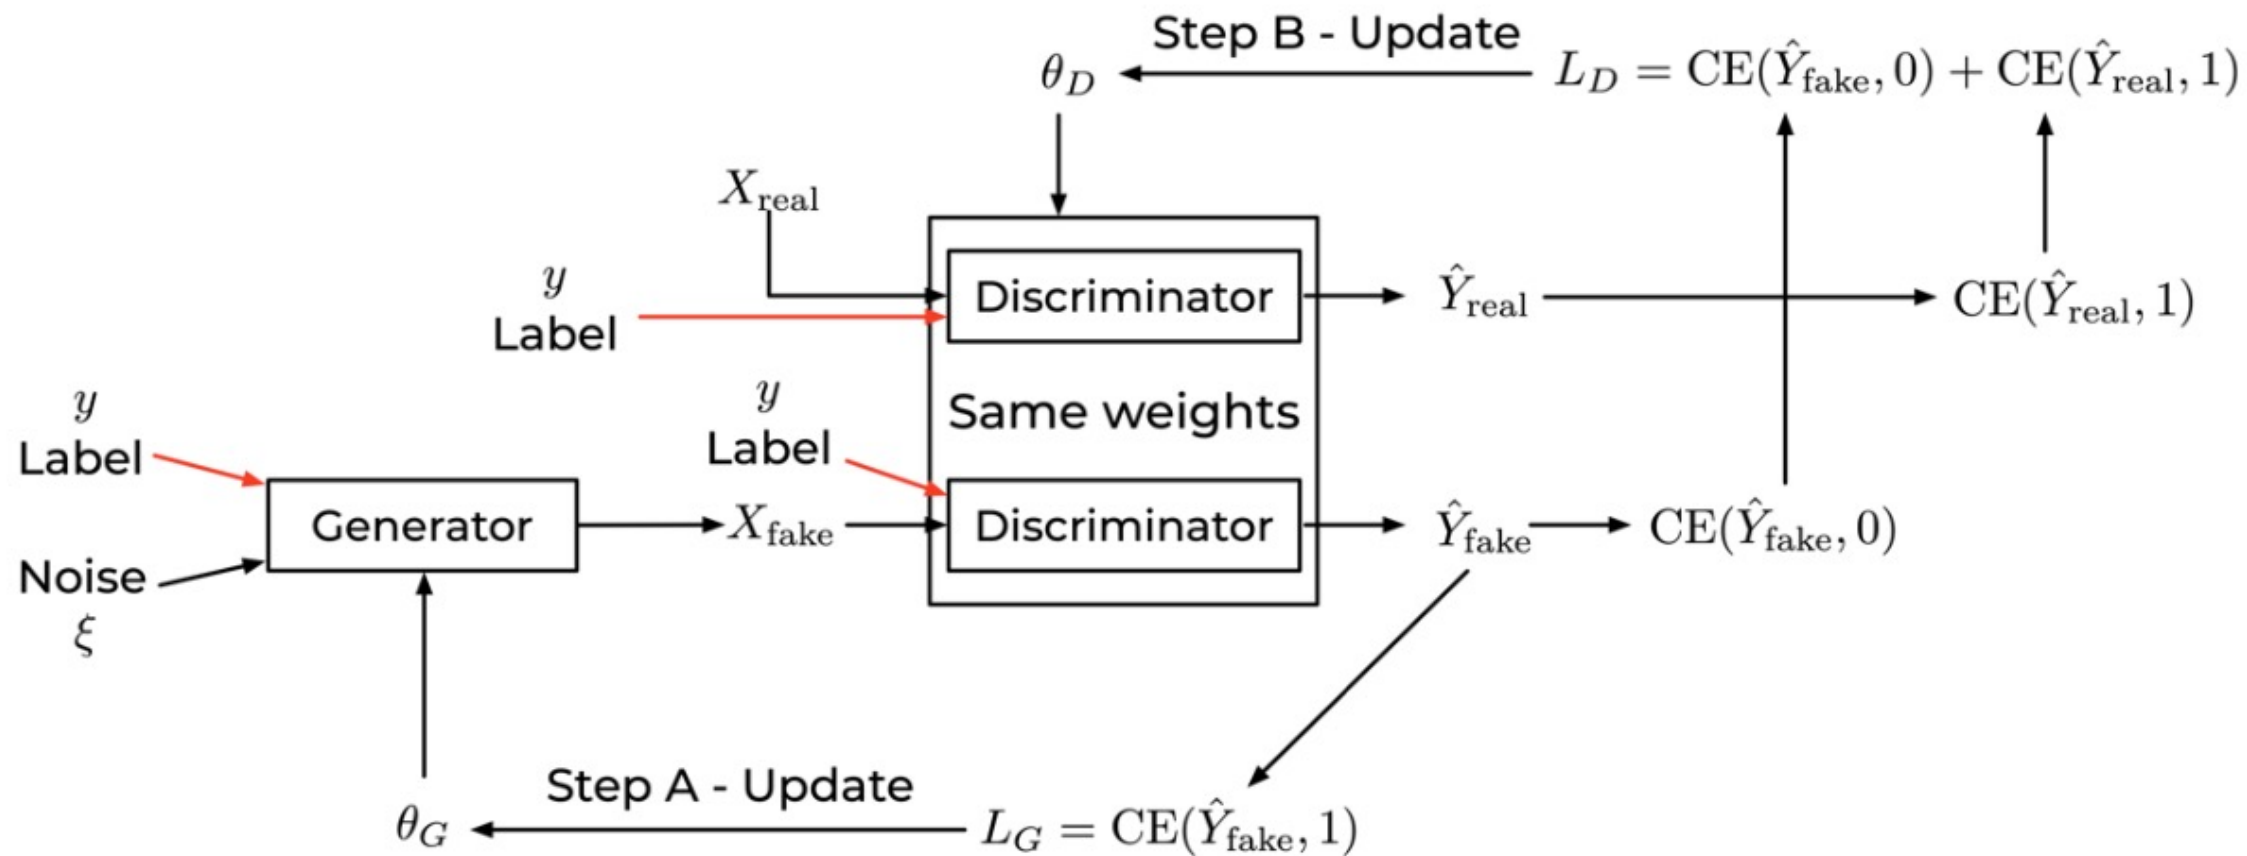
The digits were not written by a person...



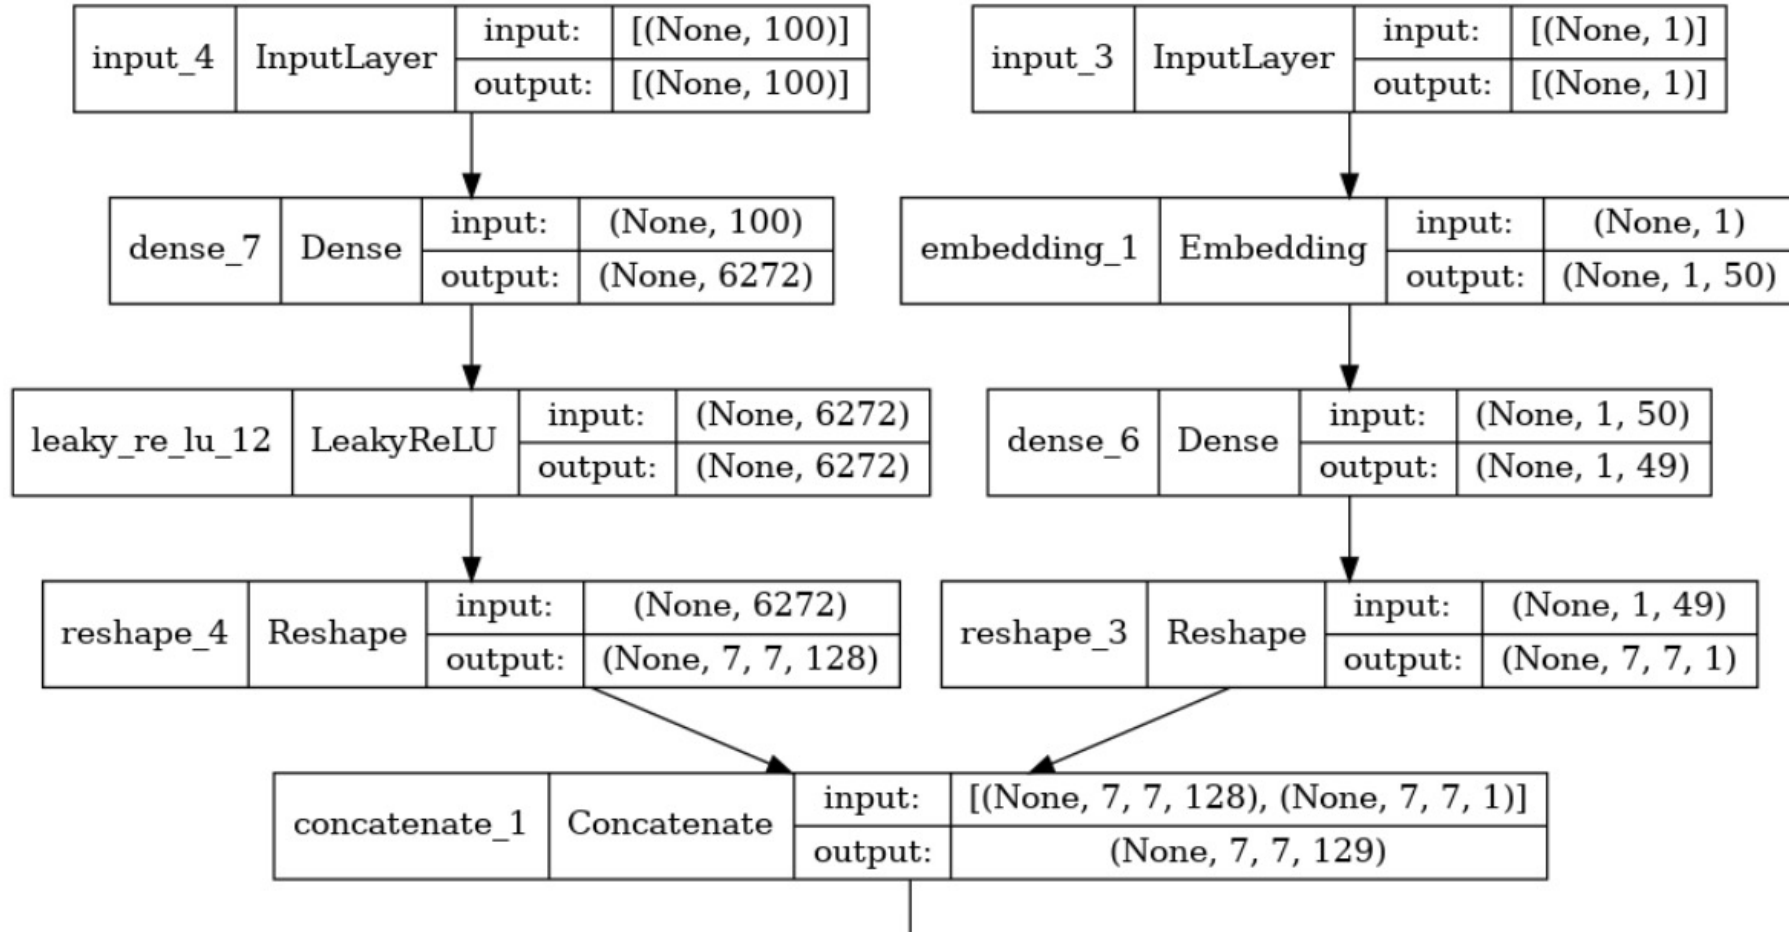
Conditional GAN

Conditional GAN (CGAN) is a **GAN variant in which both the Generator and the Discriminator are conditioned on auxiliary data such as a class label during training.**

Conditional GAN



First Layers of the generator (CGAN)



First Layers of the discriminator (CGAN)

